

Benchmarking Design document

Proposal

The objective of this assignment was to delve in the world of benchmarking the four different parts viz. CPU, memory, disk and network with an aim to understand the intricacies and have firsthand feel of the particulars, involved with getting these implemented in the real world.

Design

The overall approach for this design was to create a benchmarking module for each of the four different areas namely, CPU, Memory, Disk and Network. So, naturally this design was divided in to four parts which are mentioned further in the document.

I. CPU Benchmarking

The problem statement for this module of the benchmarking suite required stressing the CPU's FPU unit to get its throughput in terms of operations per second. This meant working with high workload which in this scenario was 1 trillion arithmetic operations to be performed and benchmarks to be calculated for the same. Since we are working with a large payload, POSIX threads were used to cater for division of work to numerous workers in a master-worker analogy.

The computation of Operations per second using a fixed number of iterations was done by using arithmetic operations for addition and subtraction in a loop, which was repeated for 1 trillion times. The total time taken then for all the operations performed in the loop is then used to compute the Giga Operation per second as per below formula:

$$\text{Giga Ops/sec} = \frac{(\text{total operations done in the loop})}{(\text{total time taken}) * 1000000000}$$

To calculate this computation for Single precision – integer operations per second and double precision – floating point operations per second didn't required any special handling but the same was not true for half precision – short operations per second and quarter precision – character operations per second.

Since the counters used to keep track of the iterations performed are using Integer which is of single precision hence, the total time for half precision and quarter precision also included them. To counter for this the time taken by the for loop using integer was extracted in a separate run which was then deducted from the total time calculations for half and quarter precisions before computing the Giga Ops/sec using the above formula.

Since the problem statement advised for strong scaling hence a structure with variables like begin and finish and thread ID were created to track the division of work to the threads. An array of this structure was then created to store these values as and when the division of work is done in the master thread.

II. Memory Benchmarking

In this module, the benchmarking suite required testing the throughput and latency of the RAM in GB/sec and microseconds respectively. The workload for which was 1 GB of data to be read and

written, 100 times for different block sizes. Since we are again working with a large payload, POSIX threads were used to cater for division of work to numerous workers in a master-worker analogy.

As the problem statement advised for strong scaling hence the division of work to the threads is being tracked inside the function using finish variable. As read and write access for the memory were to be done sequentially and randomly, multiple memcpy and memset operations were done with the varying block sizes of 1B, 1KB, 1MB and 10MB between 2 blocks called 'workloadsizeblock' and 'blksizeblock'. Since we also had to compute the latency, by doing 100 Million operations for 1 Byte of data, hence changes were done to the for loops such that loop counter changes to 100 Million when block size of 1KB was provided as input to the program. These operations were then timed, and the total time calculated and stored to be used in the computations for throughput and latency.

Finally, the computations for calculating the throughput and latency were done by using the below formulas.

$$\text{Thoroughput in GB/sec} = \frac{(\text{Total memory Read or Written to disk})}{(\text{total time taken}) * (1024 * 1024)}$$

$$\text{Latency in microseconds} = \frac{(\text{Total time taken}) * (1000)}{(\text{total memory read/written})}$$

This experiment was done using sequential and random read/write patterns, handling was required for sequential and random transfers. For sequential this, handling was done by allocating (block size) * (division of work for that thread) space to workloadsizeblock and then repeated memcpy calls from workloadsizeblock to blksizeblock with block size increments with that size of the workloadsizeblock. For random, the approach was similar with the only change that the random location was calculated by taking absolute of the difference between the work division limits assigned to that thread and random number generator spewing a number between the starting and ending limits of the work assigned to the thread.

III. Disk Benchmarking

In this module, the benchmarking suite required testing the throughput and latency of the Disk in MB/sec and milliseconds respectively. The workload for which was 10 GB of data to be read and written for different block sizes. Since we are again working with a large payload, here too POSIX threads were used to cater for division of work to numerous workers in a master-worker analogy.

Since this benchmarking is very similar to memory benchmark, to handle for strong scaling variable like finish is created to track the division of work to the threads. Like Memory benchmarking, here too, we had to do read and writes which need to be done sequentially and randomly, this was done by using functions open, read, write, lseek, etc. for varying block sizes of 1B, 1KB, 1MB and 10MB between 2 blocks called 'workloadsizeblock' and 'blksizeblock'. Since we also had to compute the latency, by doing 1 Million operations for 1 Byte of data for 1 GB workload, hence changes were done to the for loops such that loop counter changes to 1 Million when block size of 1B was provided as

input to the program. These operations were then timed, and the total time calculated and stored to be used in the computations for throughput and latency.

Finally, the computations for calculating the throughput and latency were done by using the below formulas.

$$\text{Thoroughput in MB/sec} = \frac{(\text{Total memory transferred})}{(\text{total time taken}) * (1024 * 1024 * 1024)}$$

$$\text{Latency in microseconds} = \frac{(\text{Total time taken}) * (1000000)}{(\text{total memory transferred})}$$

This experiment was done using sequential and random read/write patterns, handling was required for sequential and random transfers. For sequential this, handling was done by allocating (block size) * (division of work for that thread) space to workloadsizeblock and then repeated memcpy calls from workloadsizeblock to blksizeblock with block size increments with that size of the workloadsizeblock. For random, the approach was similar with the only change that the random location was calculated by taking absolute of the difference between the work division limits assigned to that thread and random number generator spewing a number between the starting and ending limits of the work assigned to the thread.

Apart from the above, since we had to ensure that direct input and output were to be used so while doing file write and read, I have used O_SYNC to do synchronous I/O to guarantee that the call does not return before all the data has been transferred to disk thereby giving as much throughput from disk as possible, within the given framework.

IV. Network Benchmarking

In this module, the benchmarking suite required testing the throughput and latency of the Network in MB/sec and milliseconds respectively. The workload for which was 1 GB of data to be read and written for different block sizes. Since we are again working with a large payload, here too POSIX threads were used to cater for division of work to numerous workers in a master-worker analogy.

Two designs were needed for the protocols UDP and TCP working on 2 instances in a client server fashion. To emulate the protocols the client and server constantly had to send data in record sizes of 1KB and 32KB the using send and recv functions in C. In order to build connections between the server and client sockets were implemented, along with getting addresses from the hostname provided by the batch script. Special checks were put in place to ensure the entire workload is sent between the server and client.

Since we also had to compute the latency, by doing 1 Million operations for 1 Byte of data, hence changes were done to the for loops such that loop counter changes to 1 Million when block size of 1B was provided as input to the program. These operations were then timed, and the total time calculated and stored to be used in the computations for throughput and latency.

Finally, the computations for calculating the throughput and latency were done by using the below formulas.

$$\text{Thoroughput in MB/sec} = \frac{(\text{Total bytes transferred})}{(\text{total time taken}) * (1024 * 1024)}$$

$$\text{Latency in microseconds} = \frac{(\text{Total time taken}) * (1000)}{(\text{total bytes transferred})}$$

Generic Coding highlights

- Each of the benchmark module has 2 functions, one containing the actual Benchmarking logic and the other function which is used to create and monitor threads and calls the first function.
- To generate the output for all the different testcases as summarized in the assignment problem document, we also designed bash scripts to automate the submission of batch jobs for all the modules of the benchmarking suite.
- Time duration is computed using 'struct timeval' structure as it gives the duration in seconds and microseconds.

Trade-offs made:

Generic

1. The bash scripts created to submit batch jobs are sub-optimal because Hyperion and Prometheus both had limits of 10 jobs submission per user as permissible limit at any given time.
2. The program currently isn't interactive and doesn't cater to the variable inputs from the user, although an input can be made indirectly with specific parameters by modifying existing input slurm and script files, it isn't user-friendly and complex to do so.

CPU

1. Linear equations could have been used instead of arithmetic operations alone for benchmarking the CPU.
2. Due to limited understanding on how LINPACK works, the current program doesn't have the gains as seen in LINPACK thus inferior in performance when compared with the standard benchmark.

Memory

1. Although, we are using mem* functions to do read and writes on the memory, the program doesn't handle the caching issue due to which the performance shown is not pure memory throughput.
2. Limited knowledge on handling cache issues is also one of the reason the performance is sub-optimal.

Disk

1. In this too, caching issues limit the actual performance that could have been shown by the program.
2. In this program, as the caches were not cleared hence it is a mixture of disk and memory benchmarking.

Improvements that could be done:

- Improved performance can be achieved by understanding LINPACK's internal logic and by implementing AVX and FMA instructions in the code.
- Implementing dropping of caches should be done to get better performance from the disk benchmark. A possible way was to get root access and write 3 in proc fs interface `'/proc/sys/vm/drop_caches'` to drop both inode and page caches.

Performance report

This part of the document presents the performance evaluation of the benchmark suite.

All the performance metrics were obtained from Hyperion cluster.


1. CPU Benchmark:

This benchmark was done on Hyperion cluster which has the below specifications:

Model Name	Intel(R) Xeon(R) CPU E5-2670 v3
CPU MHz	2.30 GHz
Cache Size	256.00 KiB
CPU Cores (vCore)	1

To show case the performance of the benchmarks below visuals are used along with attached data.

The below graph visually showcases the gradual increase in the performance of the benchmarks, as we go from threads 1 to 4, with the lower precisions like Quarter and half showcasing a quadratic increase in the Giga operations per second given by them.




Furthermore, the almost linear increase in the Double precision showcases that it is the least impacted by the increase in number of threads. I would say this is partly due to the fact that at lower thread count

the overhead of working with lower precision is more which increases subtly but soon shows quadratic behavior as now it breaks even with the overheads mentioned previously.

So, in this scenario, the case of having increased efficiency and throughputs due to increase in number of threads is seen in the lower precisions like Quarter & Half precisions, whereas the opposite is true for Single and Double precision as they are now hitting the maximum the CPU's FPU can handle.

Below graph showcases the performance of CPU benchmark designed in this assignment against LINPACK, and Theoretical values for Double precision. Obviously, the designed benchmark is low when compared to the other two values.


This is due to normal arithmetic operations of addition and multiplication are by themselves not enough to actually stress the CPU's Floating-point unit to give its maximum performance. Using AVX, FMA instructions while solving matrix multiplications will surely give the throughput as expected in the Linpack benchmark.



The above chart also showcases how close Linpack actual performance is when compared with the theoretical values.

When, I ran the Linpack benchmarks using thread concurrency at 1,2 and 4 with the maximum value of 70 GFLOPS seen with 4 threads for double precision which is much closer to the theoretical value than my program. Due to the issues faced with clusters at both Prometheus and Hyperion I ran the Linpack via slurm jobs rather than the interactive compute nodes as the cluster was slow and running benchmarks on the compute node was becoming quite troublesome.

The below chart expands on the efficiency when compared with precision and concurrency of threads, showcasing how efficient the program is when dealing with multiple precision and number of threads.



Below table lists the throughput values in GigaOps/sec for the CPU benchmarks based on the experiments done with the benchmark implemented along with the standard deviation for the 3 times the experiment was repeated.

These throughput value are calculated using the below formulas mentioned in the design document.

GigaOps/sec Throughput values

Workload	Concurrency	MyCPUBenchmark Measured Ops/Sec (GigaOPS)	HPL Measure rd Ops/Sec (GigaOPS)	Theoretical Ops/Sec (GigaOPS)	MyCPU Bench Efficiency(%)	HPL Efficiency(%)	Standard Deviation
QP	1	3.520475333	NA	588.8	0.597906816	NA	0.024341295
QP	2	6.484087333	NA	588.8	1.101237659	NA	0.141475038
QP	4	6.843448	NA	588.8	1.16227038	NA	0.05388375
HP	1	1.880600667	NA	294.4	0.638790987	NA	0.05460391
HP	2	3.532577	NA	294.4	1.199924253	NA	0.070374278
HP	4	3.805881	NA	294.4	1.292758492	NA	0.026022483
SP	1	3.594858333	NA	147.2	2.442159194	NA	0.1603078
SP	2	7.010844333	NA	147.2	4.762801857	NA	0.353108397


SP	4	3.64164466 7	NA	147.2	2.4739433 88	NA	0.0685090 91
DP	1	1.36466533 3	38.1844	73.6	1.8541648 55	51.880978 26	0.0380207 27
DP	2	5.075362	73.4327	73.6	6.8958722 83	99.772690 22	0.1985044 33
DP	4	10.2622803 3	70.1092	73.6	13.943315 67	95.257065 22	0.4038458 08

2. Memory Benchmark


The Memory benchmark was also done on the Hyperion cluster which has the 2133 MHz 12GB RAM sticks of which only 4GB are available to the user or node.

The below visual aids depict the performance and trends seen in this part of the assignment.


The below graph illustrates the throughput performance of Random Read and Writes as well as Sequential Read and Writes done to the memory, concurrency and block wise. As seen below the Random read and writes showcases the best performance for 1MB block size for 4 threads.



The next graph illustrates the Efficiency of the RAM benchmarking when seen mapped against the concurrency and block sizes. Here too, the best efficiency is seen at 1 MB block size and 4 threads.



Lastly the below graph showcases the performance of the MyRAMBenchmark in terms of latency and concurrency for both Random and Sequential read writes.



Below are the outputs from the benchmarking experiments

Throughput results:

Work-Load	Concurren-cy	Block Size	MyRAMBe-nch Measured Throughput (GB/sec)	pmbw Measured Throughput(GB /sec)	Theoreti-cal Throughput (GB/sec)	MyRAMBe-nch Efficiency(%)	pmbw Efficiency (%)	Standard Deviation
RWS	1	1KB	4.379744333	16.755738	63.568354	6.889819333	26.35861566	0.08100662
RWS	2	1KB	7.696003	16.161646	63.568354	12.10665767	25.42404279	0.14406986
RWS	4	1KB	8.359899333	14.842147	63.568354	13.15103967	23.34832747	0.12667387
RWS	1	1MB	4.451564667	27.834186	63.568354	7.0028	43.78623111	0.13976822
RWS	2	1MB	4.456042333	29.459142	63.568354	10.276272	46.3424646	0.26543761
RWS	4	1MB	8.177533333	28.464898	63.568354	12.86415833	44.77840945	0.21775537
RWS	1	10MB	6.522010333	34.55994	63.568354	10.259838	54.36658039	1.17521413
RWS	2	10MB	11.13223033	34.273696	63.568354	17.512221	53.9162865	0.59632099
RWS	4	10MB	12.17623167	30.645449	63.568354	19.15454933	48.20865651	2.67423325
RW R	1	1KB	8.627884333	12.843576	63.568354	13.57260933	20.20435542	0.19058895
RW R	2	1KB	3.495027	5.4529531	63.568354	5.542485333	8.578093959	0.01193508
RW R	4	1KB	2.938921667	6.2556131	63.568354	4.623246333	9.840766152	0.09909288
RW R	1	1MB	12.22489533	9.2620952	63.568354	19.231103	14.57029263	0.17266233
RW R	2	1MB	20.84701033	25.186064	63.568354	32.794636	39.62044332	0.59780384
RW R	4	1MB	24.43574333	30.584671	63.568354	38.440107	48.1130461	0.8137159
RW R	1	10MB	11.11102567	23.172895	63.568354	16.407398	36.45350823	1.07384184
RW R	2	10MB	19.624958	32.868451	63.568354	30.87221367	51.70568225	1.65886957
RW R	4	10MB	21.83599367	34.683363	63.568354	34.35041567	54.56073853	1.27941268

Latency results:

Work-load	Concurrence	Block Size	MyRAMBench Measured Latency(us)	pmbw Measure d Latency (us)	Theoretical Latency (us)	MyRAMBench Efficiency (%)	PMBW Efficiency (%)	Standard Deviation
RWS	1	1B	0.007177333	0.006319	0.01406	51.049992	44.943101	0.00029666
RWS	2	1B	0.003247667	0.003148	0.01406	23.09752967	22.38975818	0.00019356
RWS	4	1B	0.002832	0.001874	0.01406	20.14460433	13.32859175	0.00013945
RWR	1	1B	0.140941	0.143937	0.01406	9.715076333	1023.733997	0.00563812
RWR	2	1B	0.237287333	0.067617	0.01406	16.87674933	480.9174964	0.00670277
RWR	4	1B	0.326609333	0.033241	0.01406	23.22968267	236.4224751	0.00697526

Similar to Linpack, the PMBW benchmarks were submitted by me via slurm jobs the

3. Disk Benchmark

The Memory benchmark was also done on the Hyperion cluster which has the Seagate Constellation 2 SATA Hard drives. The detailed specifications for the same can be found at the below link:

<https://www.seagate.com/files/www-content/product-content/constellation-fam/constellation/constellation-2/en-us/docs/constellation2-fips-ds1719-4-1207us.pdf>

The below table showcases the results obtained from running the benchmark on Hyperion cluster.

Throughput results:


Work-load	Concurrence	Block Size	MyDisk Bench Measured Throughput(MB/sec)	IOZone Measured Throughput (MB/sec)	Theoretical Throughput (MB/sec)	MyDisk Bench Efficiency (%)	IOZone Efficiency (%)	Standard Deviation
RR	1	1MB	53.75619	313.18	600	8.959365	52.19667	11.37474
RR	2	1MB	44.72444	353.18	600	7.454074	58.86333	37.21378
RR	4	1MB	96.50222	303.18	600	16.0837	50.53	17.50689
RR	1	10MB	68.39111	526.36	600	11.39852	87.72667	7.653445

RR	2	10MB	83.05778	406.36	600	13.84296	67.72667	41.689 33
RR	4	10MB	97.73788	566.36	600	16.28965	94.39333	342.67 06
RR	1	100M B	80.67037	452.72	600	13.44506	75.45333	187.49 2
RR	2	100M B	81.16889	332.72	600	13.52815	55.45333	342.67 06
RR	4	100M B	88.33051	532.72	600	14.72175	88.78667	42.588 38
RS	1	1MB	51.42628	310.52	600	8.571047	51.75333	1.5428 75
RS	2	1MB	77.32781	325.52	600	12.88797	54.25333	13.783 45
RS	4	1MB	148.3362	301.52	600	24.7227	50.25333	32.195 68
RS	1	10MB	27.11404	321.04	600	4.519007	53.50667	19.321 95
RS	2	10MB	25.6744	501.04	600	4.279067	83.50667	0.1210 74
RS	4	10MB	41.83216	561.04	600	6.972027	93.50667	11.181 96
RS	1	100M B	56.6192	442.08	600	9.436534	73.68	11.698 76
RS	2	100M B	41.56502	302.08	600	6.927504	50.34667	211.40 2
RS	4	100M B	81.05684	302.08	600	13.50947	50.34667	10.082 52
WR	1	1MB	40.54739	284.25	600	6.757899	47.375	0.5208 62
WR	2	1MB	48.39925	219.25	600	8.066541	36.54167	0.2858 4
WR	4	1MB	62.19073	184.25	600	10.36512	30.70833	3.1185 32
WR	1	10MB	53.3278	598.5	600	8.887967	99.75	4.1336 56
WR	2	10MB	48.31124	468.5	600	8.051874	78.08333	4.1336 56
WR	4	10MB	57.36965	588.5	600	9.561609	98.08333	4.1336 56
WR	1	100M B	50.65512	426	600	8.442521	71	1.0501 17
WR	2	100M B	55.22669	367	600	9.204448	61.16667	8.0905 61
WR	4	100M B	70.09879	483	600	11.68313	80.5	23.207 78


WS	1	1MB	50.21623	185.63	600	8.369371	30.93833	1.083086
WS	2	1MB	63.03881	195.63	600	10.50647	32.605	3.951009
WS	4	1MB	87.61114	105.63	600	14.60186	17.605	1.202421
WS	1	10MB	40.25868	371.26	600	6.70978	61.87667	3.214788
WS	2	10MB	49.0156	351.26	600	8.169267	58.54333	1.020333
WS	4	10MB	58.10763	251.26	600	9.684605	41.87667	1.591361
WS	1	100MB	67.33454	410.469	600	11.22242	68.4115	3.449329
WS	2	100MB	118.4261	562.52	600	19.73768	93.75333	15.78488
WS	4	100MB	88.23094	540.458	600	14.70516	90.07633	3.317427

The below visual aids showcase the performance of the Disk benchmark implemented in the assignment.


The below graph showcases the performance of the benchmark concurrency wise, implying the obvious that the parallelization of the work for read and write improves the performance as shown by the blocks for group of 4 threads in the below chart.



This is further enhanced when comparing the efficiency as seen in the graph below. This graph also showcases the same result that the increase in the number of threads



The below graph, illustrates the latency performance for the current benchmark for varying number of threads, the data for which was taken from the table mentioned below:



This is calculated using the below formula:

Theoretical value calculation:

$$IOPS = \frac{1}{\text{average seek time} + \text{average latency}}$$

Here for the specification in Hyperion the calculation using the values is as follows:

$$\text{average seek time (read/write)} = \frac{0.0085}{0.0095} \text{ms}$$

$$\text{average latency} = 4.16 \text{ms}$$

This gives us IOPS as, $IOPS = 1/(\text{average seek time} + \text{average latency})$

So we have,

$$IOPS = 5782.4991$$

Latency values:

Work-load	Con-curren-cy	Block Size	MyDiskBench Measured Latency(ms)	IOZone Measure d Latency(ms)	Theoretic al Latency(ms)	MyDiskBench Efficiency(%)	IOZone Efficiency(%)	StandardDevia-tion
RR	1	1KB	0.0124	0.0217	4.16	0.002981	0.528846	0.001642
	2	1KB	0.0144	0.0273	4.16	0.003462	0.663462	0.000945
	4	1KB	0.00722	0.0347	4.16	0.001736	0.841346	0.000357
	8	1KB	0.00298	0.0264	4.16	0.000716	0.641827	0.0002
	16	1KB	0.00167	0.0067	4.16	0.000401	0.168269	0.00006
	32	1KB	0.00083	0.0357	4.16	0.0002	0.865385	5.77E-05
	64	1KB	0.00068	0.0167	4.16	0.000163	0.408654	5.78E-05
	128	1KB	0.00037	0.0777	4.16	8.89E-05	1.875	0.002136
WR	1	1KB	0.0143	0.1557	4.16	0.003438	3.75	0.021414
	2	1KB	0.01018	0.0377	4.16	0.002447	0.913462	0.00089
	4	1KB	0.004788	0.3637	4.16	0.001151	8.75	0.002906
	8	1KB	0.002934	0.0257	4.16	0.000705	0.625	0.001378
	16	1KB	0.001388	0.0453	4.16	0.000334	1.096154	0.000524
	32	1KB	0.008667	0.1057	4.16	0.002083	2.548077	5.83E-05
	64	1KB	0.006987	0.0553	4.16	0.00168	1.336538	5.09E-05
	128	1KB	0.005454	0.0406	4.16	0.001311	0.983173	0.000151

Apart from the above, the IOPS values were also calculated and are presented as below:

Work-load	Con-curren-cy	Block Size	MyDisk Bench Measure d IOPS	IOZone Measure d IOPS	Theoretic al IOPS	MyDiskBench Efficiency(%)	IOZone Efficiency(%)	Standar d Deviation
RR	1	1KB	4466.33	5065.118	5782.499	77.23875	87.59392	11.8781
RR	2	1KB	4627.106	5241.118	5782.499	80.01914	90.63759	7.985152
RR	4	1KB	919.725	3432.645	5782.499	15.90532	59.36265	11.60871

RR	8	1KB	2838.59 6	3332.57 8	5782.499	49.08943	57.63214	28.8272
RR	16	1KB	4403.48 5	4639.17 5	5782.499	76.15193	80.22785	58.5121 4
RR	32	1KB	3993.69 9	4608.38 9	5782.499	69.06528	79.69545	54.6959 7
RR	64	1KB	1224.62 6	2308.11 6	5782.499	21.17814	39.91555	143.801 8
RR	128	1KB	1740.59 2	3730.64 2	5782.499	30.10103	64.51609	3.56E- 13
WR	1	1KB	331.026 2	1330.71 6	5782.499	5.724622	23.01282	0.65923 7
WR	2	1KB	602.313 4	3302.00 3	5782.499	10.41614	57.1034	0.12013 1
WR	4	1KB	1423.92 2	2223.56 7	5782.499	24.62468	38.45339	0.71548 2
WR	8	1KB	2952.67 6	4310.37 5	5782.499	51.06229	74.54174	1.72680 7
WR	16	1KB	3916.50 5	4716.17 6	5782.499	67.73031	81.55947	4.81055 6
WR	32	1KB	793.890 1	3163.35 5	5782.499	13.72919	54.70568	6.00063 1
WR	64	1KB	1051.23 3	1989.00 6	5782.499	18.17956	34.39699	8.23389 5
WR	128	1KB	1354.59 3	2430.58 8	5782.499	23.42573	42.03352	16.1955 7

IOZONE was the benchmark for benchmarking Disk below are some screenshots of the same:

```

dkaramchandani@compute-3:/tmp
Alexey Skidanov.

Run began: Tue Mar 27 13:59:06 2018

File size set to 10485760 kB
Record Size 1024 kB
SYNC Mode.
Include fsync in write timing
O_DIRECT feature enabled
Command line used: /usr/bin/iozone -s 10g -r 1m -o -l2 -u2 -i 0 -i 1 -i 2 -e -I -F -f /tmp/iozonetmp1_d.txt /tmp/iozonetmp2_d.txt
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
Min process = 2
Max process = 2
Throughput test with 2 processes
Each process writes a 10485760 kByte file in 1024 kByte records

Children see throughput for 2 initial writers = 284463.17 kB/sec
Parent sees throughput for 2 initial writers = 284421.49 kB/sec
Min throughput per process = 142207.39 kB/sec
Max throughput per process = 142255.78 kB/sec
Avg throughput per process = 142231.59 kB/sec
Min xfer = 10482688.00 kB

Children see throughput for 2 rewriters = 310301.44 kB/sec
Parent sees throughput for 2 rewriters = 310283.04 kB/sec
Min throughput per process = 155014.59 kB/sec
Max throughput per process = 155286.84 kB/sec
Avg throughput per process = 155150.72 kB/sec
Min xfer = 10468352.00 kB

Children see throughput for 2 readers = 590382.09 kB/sec
Parent sees throughput for 2 readers = 590318.54 kB/sec
Min throughput per process = 292360.03 kB/sec
Max throughput per process = 298022.06 kB/sec
Avg throughput per process = 295191.05 kB/sec
Min xfer = 10286080.00 kB

Children see throughput for 2 re-readers = 590659.00 kB/sec
Parent sees throughput for 2 re-readers = 590642.40 kB/sec
Min throughput per process = 293761.91 kB/sec

```

```

dkaramchandani@compute-3:/tmp

Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
Min process = 2
Max process = 2
Throughput test with 2 processes
Each process writes a 10485760 kByte file in 1024 kByte records

Children see throughput for 2 initial writers = 284463.17 kB/sec
Parent sees throughput for 2 initial writers = 284421.49 kB/sec
Min throughput per process = 142207.39 kB/sec
Max throughput per process = 142255.78 kB/sec
Avg throughput per process = 142231.59 kB/sec
Min xfer = 10482688.00 kB

Children see throughput for 2 rewriters = 310301.44 kB/sec
Parent sees throughput for 2 rewriters = 310283.04 kB/sec
Min throughput per process = 155014.59 kB/sec
Max throughput per process = 155286.84 kB/sec
Avg throughput per process = 155150.72 kB/sec
Min xfer = 10468352.00 kB

Children see throughput for 2 readers = 590382.09 kB/sec
Parent sees throughput for 2 readers = 590318.54 kB/sec
Min throughput per process = 292360.03 kB/sec
Max throughput per process = 298022.06 kB/sec
Avg throughput per process = 295191.05 kB/sec
Min xfer = 10286080.00 kB

Children see throughput for 2 re-readers = 590659.00 kB/sec
Parent sees throughput for 2 re-readers = 590642.40 kB/sec
Min throughput per process = 293761.91 kB/sec
Max throughput per process = 296897.09 kB/sec
Avg throughput per process = 295329.50 kB/sec
Min xfer = 10375168.00 kB

Children see throughput for 2 random readers = 506618.88 kB/sec
Parent sees throughput for 2 random readers = 506590.59 kB/sec
Min throughput per process = 253299.06 kB/sec
Max throughput per process = 253320.81 kB/sec
Avg throughput per process = 253309.44 kB/sec
Min xfer = 10484736.00 kB

```

```

dkaramchandani@compute-3: /tmp
Min throughput per process = 142207.39 kB/sec
Max throughput per process = 142255.78 kB/sec
Avg throughput per process = 142231.59 kB/sec
Min xfer = 10482688.00 kB

Children see throughput for 2 rewriters = 310301.44 kB/sec
Parent sees throughput for 2 rewriters = 310283.04 kB/sec
Min throughput per process = 155014.59 kB/sec
Max throughput per process = 155286.84 kB/sec
Avg throughput per process = 155150.72 kB/sec
Min xfer = 10468352.00 kB

Children see throughput for 2 readers = 590382.09 kB/sec
Parent sees throughput for 2 readers = 590318.54 kB/sec
Min throughput per process = 292360.03 kB/sec
Max throughput per process = 298022.06 kB/sec
Avg throughput per process = 295191.05 kB/sec
Min xfer = 10286080.00 kB

Children see throughput for 2 re-readers = 590659.00 kB/sec
Parent sees throughput for 2 re-readers = 590642.40 kB/sec
Min throughput per process = 293761.91 kB/sec
Max throughput per process = 296897.09 kB/sec
Avg throughput per process = 295329.50 kB/sec
Min xfer = 10375168.00 kB

Children see throughput for 2 random readers = 506618.88 kB/sec
Parent sees throughput for 2 random readers = 506590.59 kB/sec
Min throughput per process = 253298.06 kB/sec
Max throughput per process = 253320.81 kB/sec
Avg throughput per process = 253309.44 kB/sec
Min xfer = 10484736.00 kB

Children see throughput for 2 random writers = 283166.94 kB/sec
Parent sees throughput for 2 random writers = 283140.80 kB/sec
Min throughput per process = 141574.30 kB/sec
Max throughput per process = 141592.64 kB/sec
Avg throughput per process = 141583.47 kB/sec
Min xfer = 10484736.00 kB

iozone test complete.
dkaramchandani@compute-3:/tmp$

```

Network:

I wasn't able to fully implement the socket programming bits for TCP and UDP protocols mainly due to the unstable cluster performance all over the past couple of days. The situation was really bad when the supporting cluster (Prometheus in this case) also failed to curb the growing outages faced by the students.

So, I haven't attached my network codes but I did get the benchmark values for TCP and UDP using Iperf3 some of the screenshots for the same are as below:

Iperf3 screenshots:

1 thread

```

dkaramchandani@hyperionides:~$ iperf3 -c 192.168.9.67 -P 1
Connecting to host 192.168.9.67, port 5201
[ 4] local 192.168.27.155 port 58320 connected to 192.168.9.67 port 5201
[ ID] Interval            Transfer    Bandwidth    Retr  Cwnd
[ 4] 0.00-1.00 sec        165 MBytes  1.39 Gbits/sec  41   732 KBytes
[ 4] 1.00-2.00 sec        175 MBytes  1.47 Gbits/sec  12   728 KBytes
[ 4] 2.00-3.00 sec        176 MBytes  1.48 Gbits/sec   8   707 KBytes
[ 4] 3.00-4.00 sec        176 MBytes  1.48 Gbits/sec   9   686 KBytes
[ 4] 4.00-5.00 sec        175 MBytes  1.47 Gbits/sec  10   666 KBytes
[ 4] 5.00-6.00 sec        176 MBytes  1.48 Gbits/sec   7   655 KBytes
[ 4] 6.00-7.00 sec        178 MBytes  1.49 Gbits/sec  13   648 KBytes
[ 4] 7.00-8.00 sec        176 MBytes  1.48 Gbits/sec   2   635 KBytes
[ 4] 8.00-9.00 sec        175 MBytes  1.47 Gbits/sec  29   615 KBytes
[ 4] 9.00-10.00 sec       175 MBytes  1.47 Gbits/sec   8   605 KBytes
- - - - -
[ ID] Interval            Transfer    Bandwidth    Retr
[ 4] 0.00-10.00 sec      1.71 GBytes  1.47 Gbits/sec  139
[ 4] 0.00-10.00 sec      1.70 GBytes  1.46 Gbits/sec
sender
receiver

iperf Done.
dkaramchandani@hyperionides:~$

```

2 threads

```

dkaramchandani@hyperionides:~$ iperf3 -c 192.168.9.67 -P 2
Connecting to host 192.168.9.67, port 5201
[ 4] local 192.168.27.155 port 58614 connected to 192.168.9.67 port 5201
[ 6] local 192.168.27.155 port 58616 connected to 192.168.9.67 port 5201
[ ID] Interval      Transfer    Bandwidth   Retr  Cwnd
[ 4]  0.00-1.00    sec    132 MBytes  1.11 Gbits/sec    50   583 KBytes
[ 6]  0.00-1.00    sec    102 MBytes  856 Mbits/sec    10   288 KBytes
[SUM] 0.00-1.00    sec    234 MBytes  1.97 Gbits/sec    60
- - - - -
[ 4]  1.00-2.00    sec    142 MBytes  1.19 Gbits/sec    34   460 KBytes
[ 6]  1.00-2.00    sec    100 MBytes  840 Mbits/sec    18   355 KBytes
[SUM] 1.00-2.00    sec    242 MBytes  2.03 Gbits/sec    52
- - - - -
[ 4]  2.00-3.00    sec    128 MBytes  1.07 Gbits/sec    28   406 KBytes
[ 6]  2.00-3.00    sec    120 MBytes  1.01 Gbits/sec    15   335 KBytes
[SUM] 2.00-3.00    sec    247 MBytes  2.08 Gbits/sec    43
- - - - -
[ 4]  3.00-4.00    sec    120 MBytes  1.01 Gbits/sec    27   379 KBytes
[ 6]  3.00-4.00    sec    120 MBytes  1.01 Gbits/sec    12   472 KBytes
[SUM] 3.00-4.00    sec    240 MBytes  2.01 Gbits/sec    39
- - - - -
[ 4]  4.00-5.00    sec    118 MBytes  986 Mbits/sec    26   352 KBytes
[ 6]  4.00-5.00    sec    141 MBytes  1.18 Gbits/sec    49   533 KBytes
[SUM] 4.00-5.00    sec    258 MBytes  2.17 Gbits/sec    75
- - - - -
[ 4]  5.00-6.00    sec    108 MBytes  902 Mbits/sec     4   385 KBytes
[ 6]  5.00-6.00    sec    132 MBytes  1.11 Gbits/sec    38   462 KBytes
[SUM] 5.00-6.00    sec    239 MBytes  2.01 Gbits/sec    42
- - - - -
[ 4]  6.00-7.00    sec    110 MBytes  923 Mbits/sec    38   373 KBytes
[ 6]  6.00-7.00    sec    138 MBytes  1.16 Gbits/sec    12   508 KBytes
[SUM] 6.00-7.00    sec    248 MBytes  2.08 Gbits/sec    50
- - - - -
[ 4]  7.00-8.00    sec    125 MBytes  1.05 Gbits/sec     7   475 KBytes
[ 6]  7.00-8.00    sec    123 MBytes  1.03 Gbits/sec    20   451 KBytes
[SUM] 7.00-8.00    sec    248 MBytes  2.08 Gbits/sec    27
- - - - -
[ 4]  8.00-9.00    sec    112 MBytes  944 Mbits/sec    53   308 KBytes
[ 6]  8.00-9.00    sec    133 MBytes  1.12 Gbits/sec    17   390 KBytes
[SUM] 8.00-9.00    sec    246 MBytes  2.06 Gbits/sec    70
- - - - -
[ 4]  9.00-10.00   sec    116 MBytes  975 Mbits/sec     5   413 KBytes
[ 6]  9.00-10.00   sec    134 MBytes  1.13 Gbits/sec    39   462 KBytes
[SUM] 9.00-10.00   sec    251 MBytes  2.10 Gbits/sec    44

```

4 threads

```

dkaramchandani@hyperionides: ~
[SUM] 7.00-8.00 sec 311 MBytes 2.61 Gbits/sec 635
[ 4] 8.00-9.00 sec 40.8 MBytes 342 Mbits/sec 89 158 KBytes
[ 6] 8.00-9.00 sec 41.2 MBytes 346 Mbits/sec 130 96.2 KBytes
[ 8] 8.00-9.00 sec 33.8 MBytes 284 Mbits/sec 89 103 KBytes
[10] 8.00-9.00 sec 41.8 MBytes 351 Mbits/sec 34 165 KBytes
[12] 8.00-9.00 sec 32.5 MBytes 273 Mbits/sec 51 129 KBytes
[14] 8.00-9.00 sec 48.5 MBytes 407 Mbits/sec 55 158 KBytes
[16] 8.00-9.00 sec 30.9 MBytes 260 Mbits/sec 73 102 KBytes
[18] 8.00-9.00 sec 45.6 MBytes 383 Mbits/sec 59 161 KBytes
[SUM] 8.00-9.00 sec 315 MBytes 2.64 Gbits/sec 580

[ 4] 9.00-10.00 sec 45.9 MBytes 385 Mbits/sec 39 240 KBytes
[ 6] 9.00-10.00 sec 39.5 MBytes 331 Mbits/sec 119 102 KBytes
[ 8] 9.00-10.00 sec 27.5 MBytes 231 Mbits/sec 183 174 KBytes
[10] 9.00-10.00 sec 28.8 MBytes 241 Mbits/sec 80 123 KBytes
[12] 9.00-10.00 sec 28.8 MBytes 241 Mbits/sec 102 84.8 KBytes
[14] 9.00-10.00 sec 44.6 MBytes 374 Mbits/sec 87 102 KBytes
[16] 9.00-10.00 sec 35.7 MBytes 299 Mbits/sec 63 148 KBytes
[18] 9.00-10.00 sec 54.3 MBytes 456 Mbits/sec 43 136 KBytes
[SUM] 9.00-10.00 sec 305 MBytes 2.56 Gbits/sec 716

[ ID] Interval Transfer Bandwidth Retr
[ 4] 0.00-10.00 sec 413 MBytes 347 Mbits/sec 684 sender
[ 4] 0.00-10.00 sec 412 MBytes 346 Mbits/sec 684 receiver
[ 6] 0.00-10.00 sec 376 MBytes 315 Mbits/sec 877 sender
[ 6] 0.00-10.00 sec 375 MBytes 314 Mbits/sec 877 receiver
[ 8] 0.00-10.00 sec 388 MBytes 326 Mbits/sec 993 sender
[ 8] 0.00-10.00 sec 387 MBytes 325 Mbits/sec 993 receiver
[10] 0.00-10.00 sec 380 MBytes 319 Mbits/sec 816 sender
[10] 0.00-10.00 sec 379 MBytes 318 Mbits/sec 816 receiver
[12] 0.00-10.00 sec 361 MBytes 303 Mbits/sec 988 sender
[12] 0.00-10.00 sec 360 MBytes 302 Mbits/sec 988 receiver
[14] 0.00-10.00 sec 371 MBytes 311 Mbits/sec 1019 sender
[14] 0.00-10.00 sec 370 MBytes 310 Mbits/sec 1019 receiver
[16] 0.00-10.00 sec 363 MBytes 305 Mbits/sec 835 sender
[16] 0.00-10.00 sec 363 MBytes 304 Mbits/sec 835 receiver
[18] 0.00-10.00 sec 421 MBytes 354 Mbits/sec 668 sender
[18] 0.00-10.00 sec 420 MBytes 352 Mbits/sec 668 receiver
[SUM] 0.00-10.00 sec 3.00 GBytes 2.58 Gbits/sec 6880 sender
[SUM] 0.00-10.00 sec 2.99 GBytes 2.57 Gbits/sec 6880 receiver

iperf Done.
dkaramchandani@hyperionides:~$

```

```

[ ID] Interval Transfer Bandwidth Retr
[ 4] 0.00-10.00 sec 417 MBytes 350 Mbits/sec 719 sender
[ 4] 0.00-10.00 sec 415 MBytes 348 Mbits/sec 719 receiver
[ 6] 0.00-10.00 sec 388 MBytes 326 Mbits/sec 853 sender
[ 6] 0.00-10.00 sec 387 MBytes 325 Mbits/sec 853 receiver
[ 8] 0.00-10.00 sec 337 MBytes 283 Mbits/sec 1151 sender
[ 8] 0.00-10.00 sec 336 MBytes 282 Mbits/sec 1151 receiver
[10] 0.00-10.00 sec 453 MBytes 380 Mbits/sec 495 sender
[10] 0.00-10.00 sec 452 MBytes 379 Mbits/sec 495 receiver
[12] 0.00-10.00 sec 315 MBytes 265 Mbits/sec 1167 sender
[12] 0.00-10.00 sec 314 MBytes 264 Mbits/sec 1167 receiver
[14] 0.00-10.00 sec 409 MBytes 343 Mbits/sec 773 sender
[14] 0.00-10.00 sec 408 MBytes 342 Mbits/sec 773 receiver
[16] 0.00-10.00 sec 379 MBytes 318 Mbits/sec 823 sender
[16] 0.00-10.00 sec 378 MBytes 317 Mbits/sec 823 receiver
[18] 0.00-10.00 sec 326 MBytes 273 Mbits/sec 979 sender
[18] 0.00-10.00 sec 325 MBytes 272 Mbits/sec 979 receiver
[SUM] 0.00-10.00 sec 2.95 GBytes 2.54 Gbits/sec 6960 sender
[SUM] 0.00-10.00 sec 2.94 GBytes 2.53 Gbits/sec 6960 receiver

iperf Done.
dkaramchandani@hyperionides:~$

```

Similarly for UDP as well, I managed to secure the below screenshots :

UDP :

1 thread

```

dkaramchandani@hyperionides:~$ iperf3 -c 192.168.9.67 -P 1 -u
Connecting to host 192.168.9.67, port 5201
[ 4] local 192.168.27.155 port 57471 connected to 192.168.9.67 port 5201
[ ID] Interval           Transfer     Bandwidth       Total Datagrams
[ 4]  0.00-1.00   sec      120 KBytes    983 Kbits/sec    15
[ 4]  1.00-2.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  2.00-3.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  3.00-4.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  4.00-5.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  5.00-6.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  6.00-7.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  7.00-8.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  8.00-9.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 4]  9.00-10.00  sec      128 KBytes    1.05 Mbits/sec   16
- - - - -
[ ID] Interval           Transfer     Bandwidth       Jitter          Lost/Total Datagrams
[ 4]  0.00-10.00  sec      1.24 MBytes    1.04 Mbits/sec   0.274 ms        0/159 (0%)
[ 4] Sent 159 datagrams

iperf Done.

```

2 threads

```

dkaramchandani@hyperionides:~$ iperf3 -c 192.168.9.67 -P 2 -u
Connecting to host 192.168.9.67, port 5201
[ 4] local 192.168.27.155 port 57346 connected to 192.168.9.67 port 5201
[ 6] local 192.168.27.155 port 60596 connected to 192.168.9.67 port 5201
[ ID] Interval           Transfer     Bandwidth       Total Datagrams
[ 4]  0.00-1.00   sec      120 KBytes    982 Kbits/sec    15
[ 6]  0.00-1.00   sec      120 KBytes    982 Kbits/sec    15
[SUM] 0.00-1.00   sec      240 KBytes    1.96 Mbits/sec   30
- - - - -
[ 4]  1.00-2.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  1.00-2.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 1.00-2.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  2.00-3.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  2.00-3.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 2.00-3.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  3.00-4.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  3.00-4.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 3.00-4.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  4.00-5.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  4.00-5.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 4.00-5.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  5.00-6.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  5.00-6.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 5.00-6.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  6.00-7.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  6.00-7.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 6.00-7.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  7.00-8.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  7.00-8.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 7.00-8.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  8.00-9.00   sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  8.00-9.00   sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 8.00-9.00   sec      256 KBytes    2.10 Mbits/sec   32
- - - - -
[ 4]  9.00-10.00  sec      128 KBytes    1.05 Mbits/sec   16
[ 6]  9.00-10.00  sec      128 KBytes    1.05 Mbits/sec   16
[SUM] 9.00-10.00  sec      256 KBytes    2.10 Mbits/sec   32

```

```

[  4]  5.00-6.00  sec  128 KBytes  1.05 Mbits/sec  16
[  6]  5.00-6.00  sec  128 KBytes  1.05 Mbits/sec  16
[SUM]  5.00-6.00  sec  256 KBytes  2.10 Mbits/sec  32
-----
[  4]  6.00-7.00  sec  128 KBytes  1.05 Mbits/sec  16
[  6]  6.00-7.00  sec  128 KBytes  1.05 Mbits/sec  16
[SUM]  6.00-7.00  sec  256 KBytes  2.10 Mbits/sec  32
-----
[  4]  7.00-8.00  sec  128 KBytes  1.05 Mbits/sec  16
[  6]  7.00-8.00  sec  128 KBytes  1.05 Mbits/sec  16
[SUM]  7.00-8.00  sec  256 KBytes  2.10 Mbits/sec  32
-----
[  4]  8.00-9.00  sec  128 KBytes  1.05 Mbits/sec  16
[  6]  8.00-9.00  sec  128 KBytes  1.05 Mbits/sec  16
[SUM]  8.00-9.00  sec  256 KBytes  2.10 Mbits/sec  32
-----
[  4]  9.00-10.00 sec  128 KBytes  1.05 Mbits/sec  16
[  6]  9.00-10.00 sec  128 KBytes  1.05 Mbits/sec  16
[SUM]  9.00-10.00 sec  256 KBytes  2.10 Mbits/sec  32
-----
[ ID] Interval      Transfer      Bandwidth      Jitter      Lost/Total  Datagrams
[  4]  0.00-10.00  sec  1.24 MBytes  1.04 Mbits/sec  0.198 ms    0/159 (0%)
[  4] Sent 159 datagrams
[  6]  0.00-10.00  sec  1.24 MBytes  1.04 Mbits/sec  0.197 ms    0/159 (0%)
[  6] Sent 159 datagrams
[SUM]  0.00-10.00  sec  2.48 MBytes  2.08 Mbits/sec  0.197 ms    0/318 (0%)

iperf Done.

```

8 threads

```

dkaramchandani@hyperionides:~$ iperf3 -c 192.168.9.67 -P 8 -u
Connecting to host 192.168.9.67, port 5201
[ 4] local 192.168.27.155 port 41043 connected to 192.168.9.67 port 5201
[ 6] local 192.168.27.155 port 51942 connected to 192.168.9.67 port 5201
[ 8] local 192.168.27.155 port 38680 connected to 192.168.9.67 port 5201
[10] local 192.168.27.155 port 52004 connected to 192.168.9.67 port 5201
[12] local 192.168.27.155 port 55070 connected to 192.168.9.67 port 5201
[14] local 192.168.27.155 port 49561 connected to 192.168.9.67 port 5201
[16] local 192.168.27.155 port 51946 connected to 192.168.9.67 port 5201
[18] local 192.168.27.155 port 49231 connected to 192.168.9.67 port 5201
[ ID] Interval            Transfer      Bandwidth    Total Datagrams
[ 4] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[ 6] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[ 8] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[10] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[12] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[14] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[16] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[18] 0.00-1.00 sec      120 KBytes   983 Kbits/sec  15
[SUM] 0.00-1.00 sec      960 KBytes   7.86 Mbits/sec  120
- - - - -
[ 4] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[ 6] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[ 8] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[10] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[12] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[14] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[16] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[18] 1.00-2.00 sec      128 KBytes   1.05 Mbits/sec  16
[SUM] 1.00-2.00 sec      1.00 MBytes  8.39 Mbits/sec  128
- - - - -
[ 4] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[ 6] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[ 8] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[10] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[12] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[14] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[16] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[18] 2.00-3.00 sec      128 KBytes   1.05 Mbits/sec  16
[SUM] 2.00-3.00 sec      1.00 MBytes  8.39 Mbits/sec  128
- - - - -
[ 4] 3.00-4.00 sec      128 KBytes   1.05 Mbits/sec  16
[ 6] 3.00-4.00 sec      128 KBytes   1.05 Mbits/sec  16
[ 8] 3.00-4.00 sec      128 KBytes   1.05 Mbits/sec  16

```



```

[ 18] 7.00-8.00 sec 128 KBytes 1.05 Mbits/sec 16
[SUM] 7.00-8.00 sec 1.00 MBytes 8.39 Mbits/sec 128
- - - - -
[ 4] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[ 6] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[ 8] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[10] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[12] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[14] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[16] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[18] 8.00-9.00 sec 128 KBytes 1.05 Mbits/sec 16
[SUM] 8.00-9.00 sec 1.00 MBytes 8.39 Mbits/sec 128
- - - - -
[ 4] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[ 6] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[ 8] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[10] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[12] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[14] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[16] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[18] 9.00-10.00 sec 128 KBytes 1.05 Mbits/sec 16
[SUM] 9.00-10.00 sec 1.00 MBytes 8.39 Mbits/sec 128
- - - - -
[ ID] Interval Transfer Bandwidth Jitter Lost/Total Datagrams
[ 4] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 1.245 ms 0/159 (0%)
[ 4] Sent 159 datagrams
[ 6] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 1.099 ms 0/159 (0%)
[ 6] Sent 159 datagrams
[ 8] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 0.949 ms 0/159 (0%)
[ 8] Sent 159 datagrams
[10] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 0.963 ms 0/159 (0%)
[10] Sent 159 datagrams
[12] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 0.982 ms 0/159 (0%)
[12] Sent 159 datagrams
[14] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 0.898 ms 0/159 (0%)
[14] Sent 159 datagrams
[16] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 0.907 ms 0/159 (0%)
[16] Sent 159 datagrams
[18] 0.00-10.00 sec 1.24 MBytes 1.04 Mbits/sec 0.791 ms 0/159 (0%)
[18] Sent 159 datagrams
[SUM] 0.00-10.00 sec 9.94 MBytes 8.34 Mbits/sec 0.979 ms 0/1272 (0%)

iperf Done.

```